



Eight principles for Cloud Native Storage

Cheryl Hung @oicheryl



Cheryl
@oicheryl



Why do I need storage?

Why do I need storage?



Why do I need storage?



App
binaries



App
data



Config



Backup

Why is this tricky with containers?



No
storage
pets

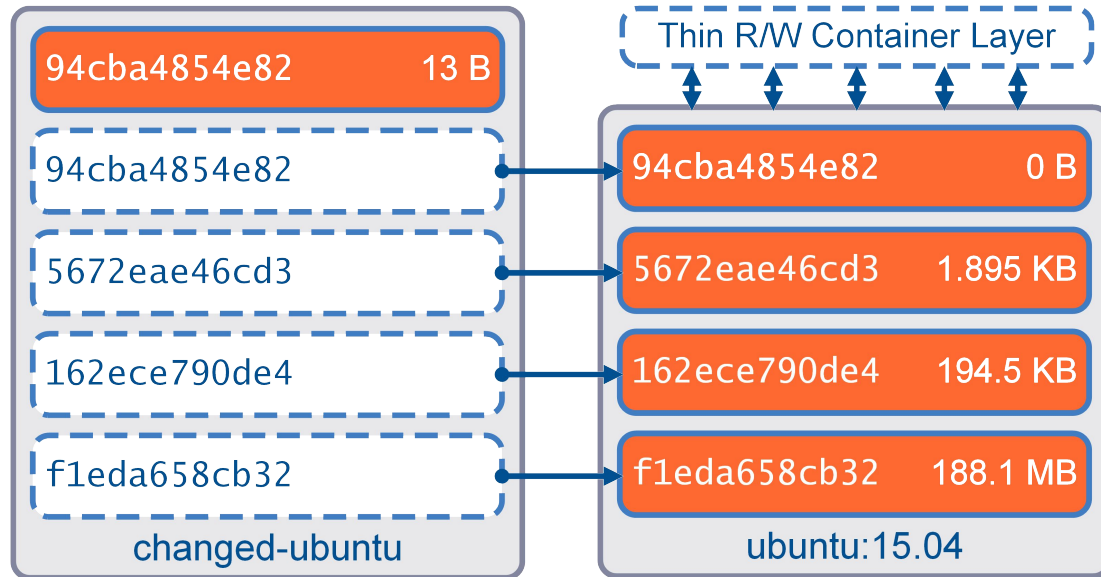


Data follows



Humans are fallible

Docker container layers




Docker local volumes

```
$ docker volume create --name mydata
```

```
$ docker run --rm -v mydata:/data:rw alpine ash -c \  
"echo hello world > /data/myfile"
```

```
$ sudo cat /var/lib/docker/volumes/mydata/_data/myfile  
hello world
```



Eight principles of Cloud Native Storage



What is Cloud Native?

Horizontally scalable

No single point of failure

Resilient and self healing

Minimal operator overhead

Decoupled from the underlying platform

Eight principles of Cloud Native Storage

1. Platform agnostic

Eight principles of Cloud Native Storage

1. Platform agnostic
- 2. API driven**

Eight principles of Cloud Native Storage

1. Platform agnostic
2. API driven
- 3. Declarative and composable**

Eight principles of Cloud Native Storage

1. Platform agnostic
2. API driven
3. Declarative and composable
- 4. Application centric**

Eight principles of Cloud Native Storage

1. Platform agnostic
2. API driven
3. Declarative and composable
4. Application centric

5. Agile

Eight principles of Cloud Native Storage

1. Platform agnostic
2. API driven
3. Declarative and composable
4. Application centric
5. Agile
- 6. Natively secure**

Eight principles of Cloud Native Storage


1. Platform agnostic
2. API driven
3. Declarative and composable
4. Application centric
5. Agile
6. Natively secure
- 7. Performant**

Eight principles of Cloud Native Storage

1. Platform agnostic
2. API driven
3. Declarative and composable
4. Application centric
5. Agile
6. Natively secure
7. Performant
- 8. Consistently available**

Storage landscape





Block storage

Data stored in fixed-size 'blocks' in a rigid arrangement—ideal for enterprise databases



File storage


Data stored as 'files' in hierarchically nested 'folders'—ideal for active documents



Object storage


Data stored as 'objects' in scalable 'buckets'—ideal for unstructured big data, analytics and archiving

Centralised file system: NFS




Centralised file system: NFS

0



Storage array: Dell EMC




Storage array: Dell EMC

Deterministic performance

Vendor lock in

2

Distributed: Ceph



Distributed: Ceph

4


Horizontally scalable

Hardware agnostic

Complicated to set up (see Rook)

Failures are expensive

Public cloud: AWS EBS



Public cloud: AWS EBS

6

Horizontally scalable


Consistent and performant

Vendor lock in


Mounting physical block devices

Expensive and privacy issues

Plugin framework: REX-Ray



Volume plugin: StorageOS



Volume plugin: StorageOS


8



Volume plugin: StorageOS



Conclusion



K8S Storage SIG & CNCF Storage WG: <https://github.com/cncf/wg-storage>

Objective is to define an industry standard “Container Storage Interface” (CSI) that will enable storage vendors (SP) to develop a plugin once and have it work across a number of container orchestration (CO) systems.



Thanks

Slides at oicheryl.com

